



Copyright And Confidentiality

All content of this material, including text, images, audio, video, software, programs, tools etc., are collected online.

Visitors may use the content or services provided in this material for personal study, research or appreciation, as well as other non-commercial or non-profit uses, but at the same time comply with the provisions of the Copyright Law and other relevant laws, and may not infringe this information and related the legal rights of the right holder. In addition, any use of any content or service of this material for other purposes requires the written permission of the person and payment.



牧云/罗伟

Network | IaaS | PaaS | ServiceMesh

交流 学习 沉淀 成长 分享

olaf.luo@foxmail.com

<https://www.yuque.com/wei.luo>

<https://youdianzhishi.com/web>

All Rights Reserved.

Rowan Luo

交流

学习

沉淀

成长

分享

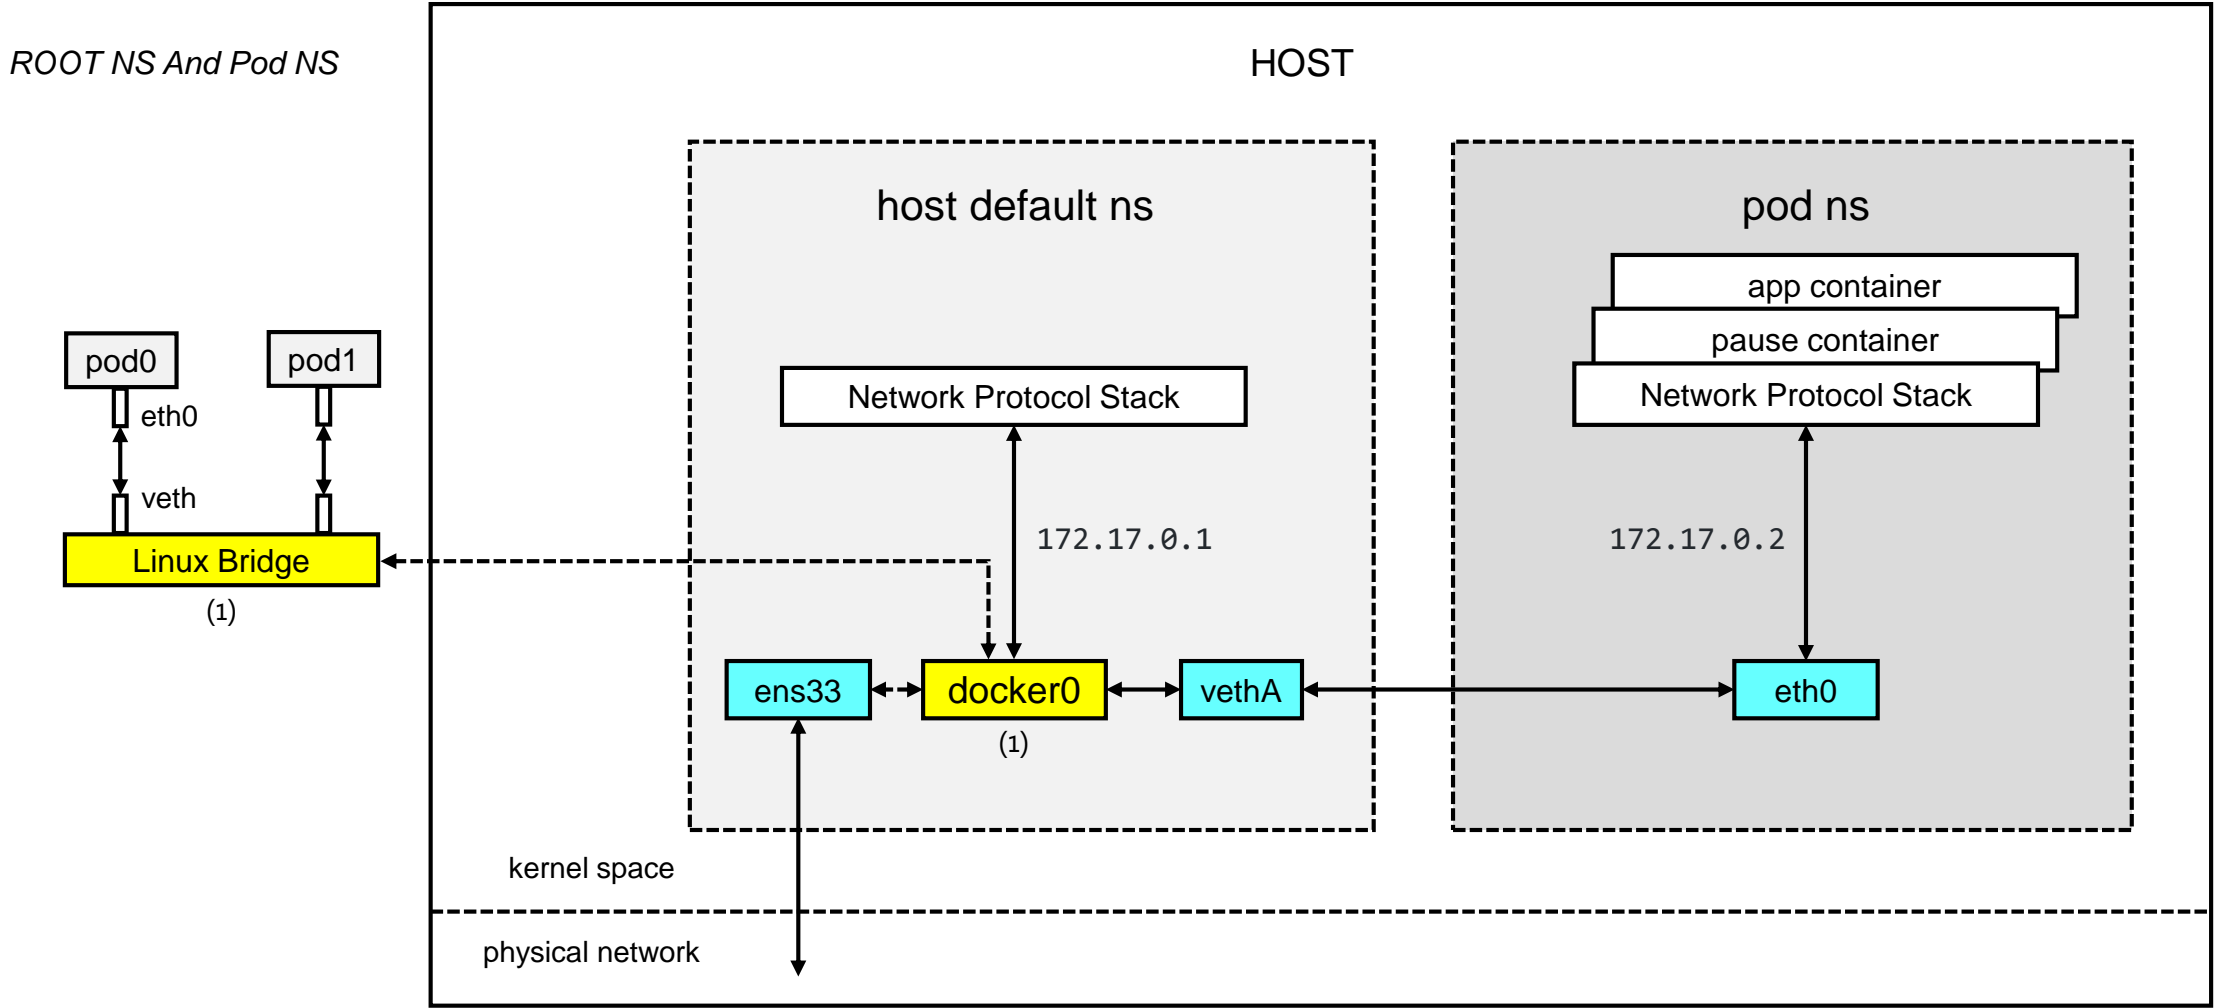
Kubernetes Network

Revision Date: 2022/02/21

Version: v1.2

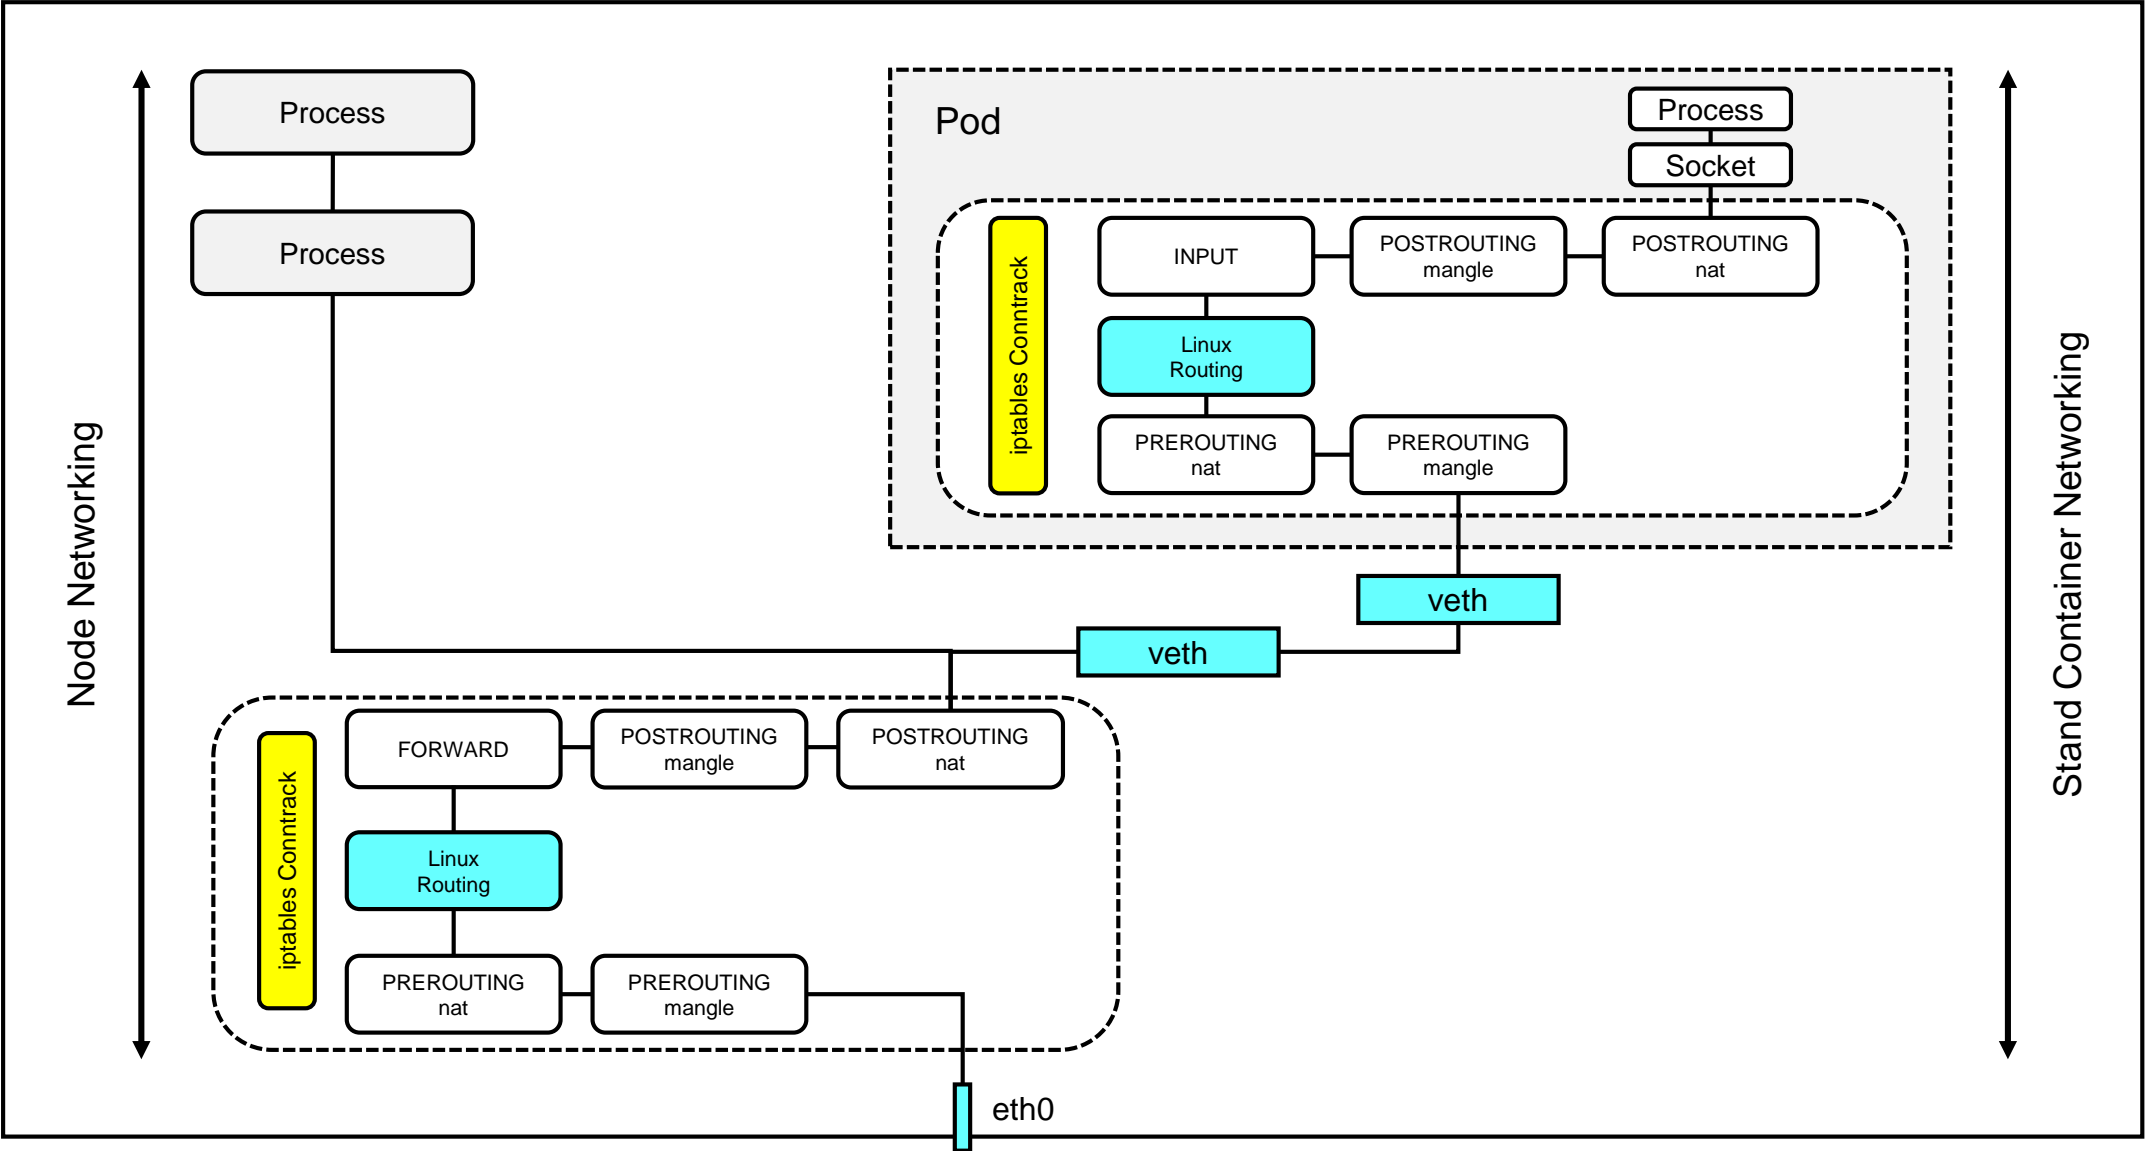
DocID: CN00002022XLW

Kubernetes CNI - Pod Network

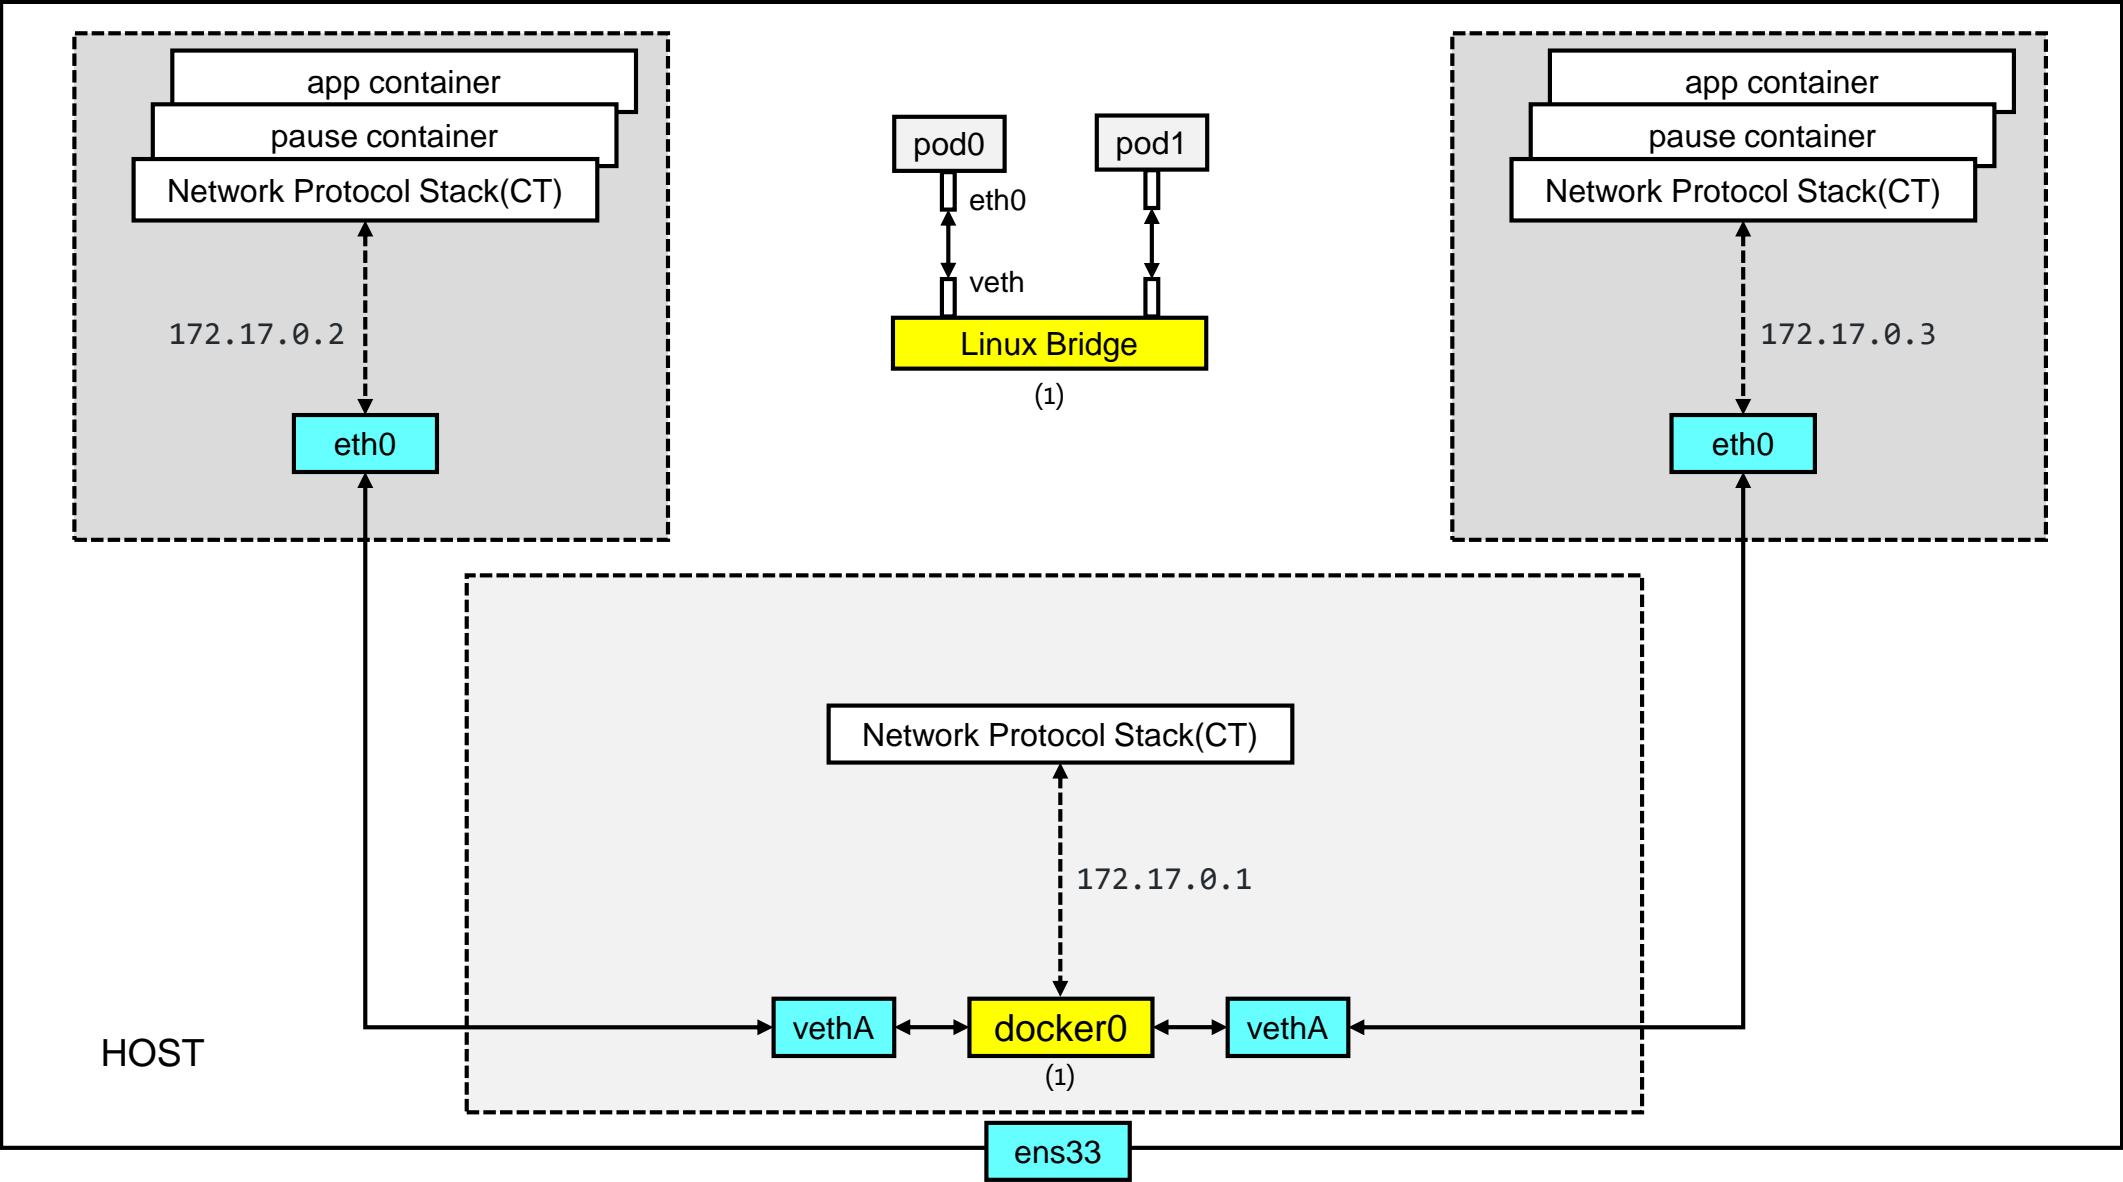


Docker spawns a container in the containers own network namespace (use the `CLONE_NEWNET` flag defined in `sched.h` when calling the `clone` system call to create a new network namespace) and later on runs a `veth` pair (a cable with two ends) between the container namespace and the host network stack.

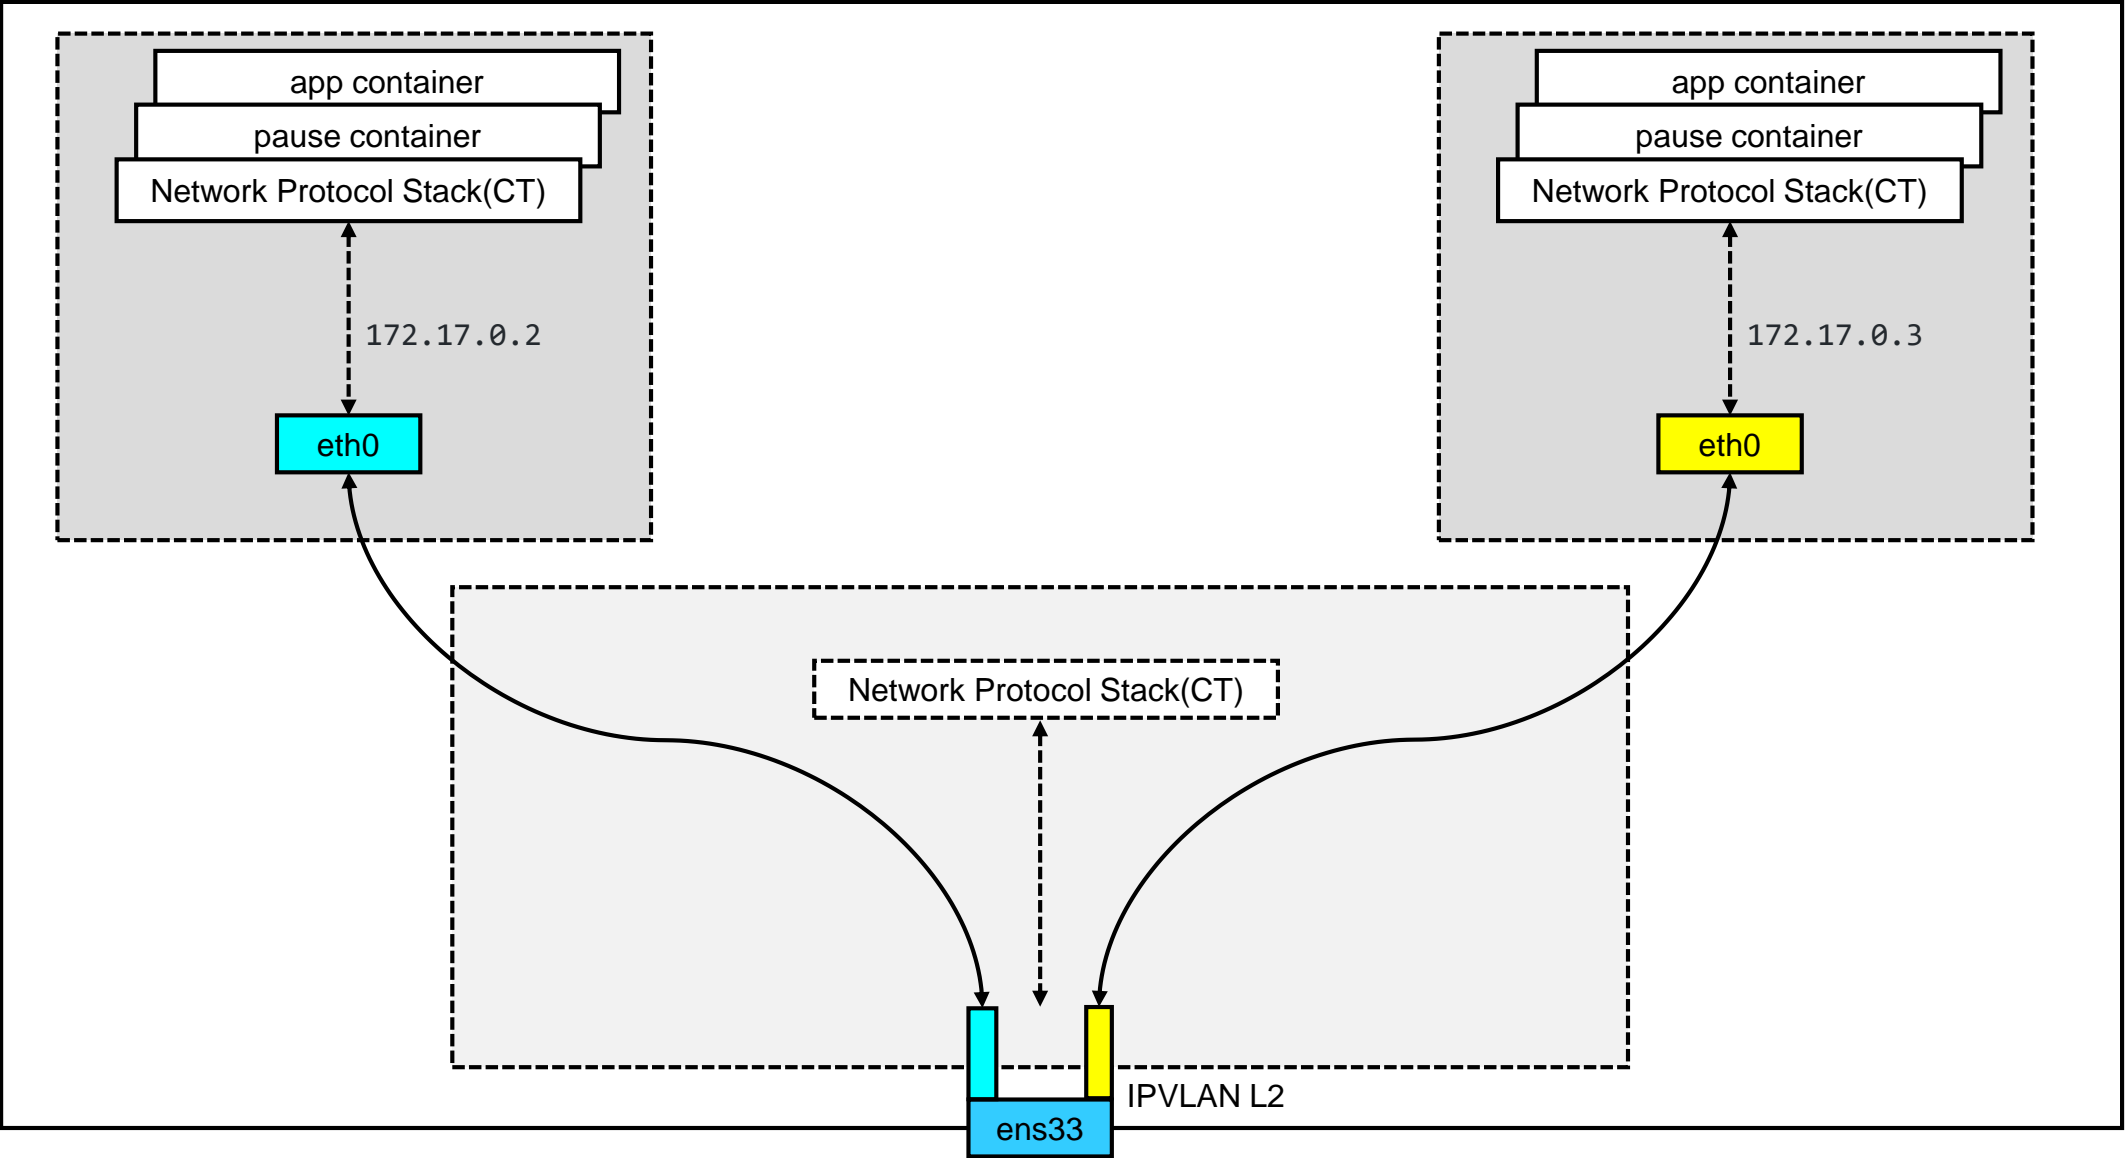
Kubernetes CNI - Pod Network vs Node Network



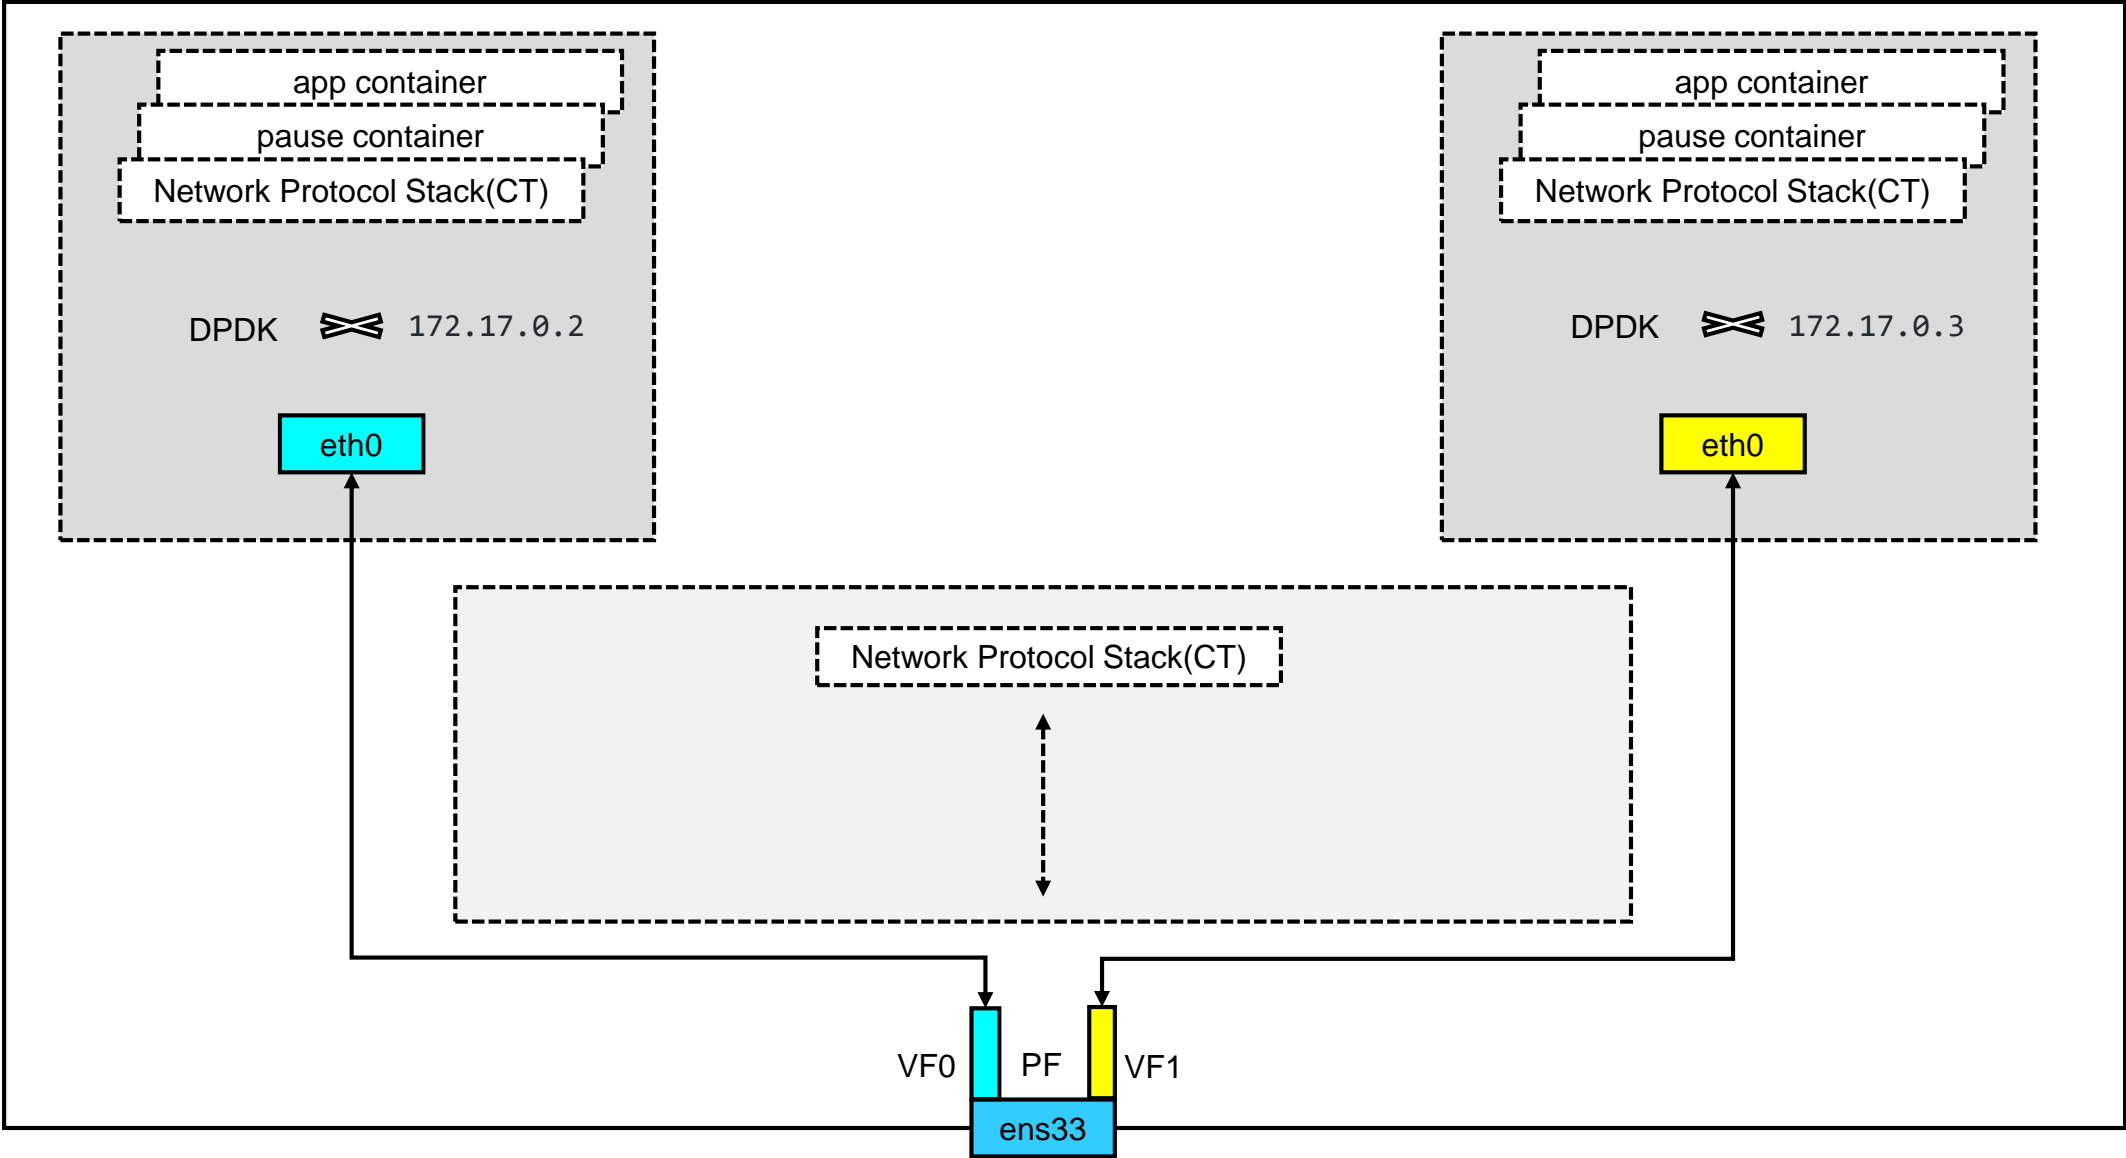
Kubernetes CNI - Bridge Mode



Kubernetes CNI - NIC Multiplex

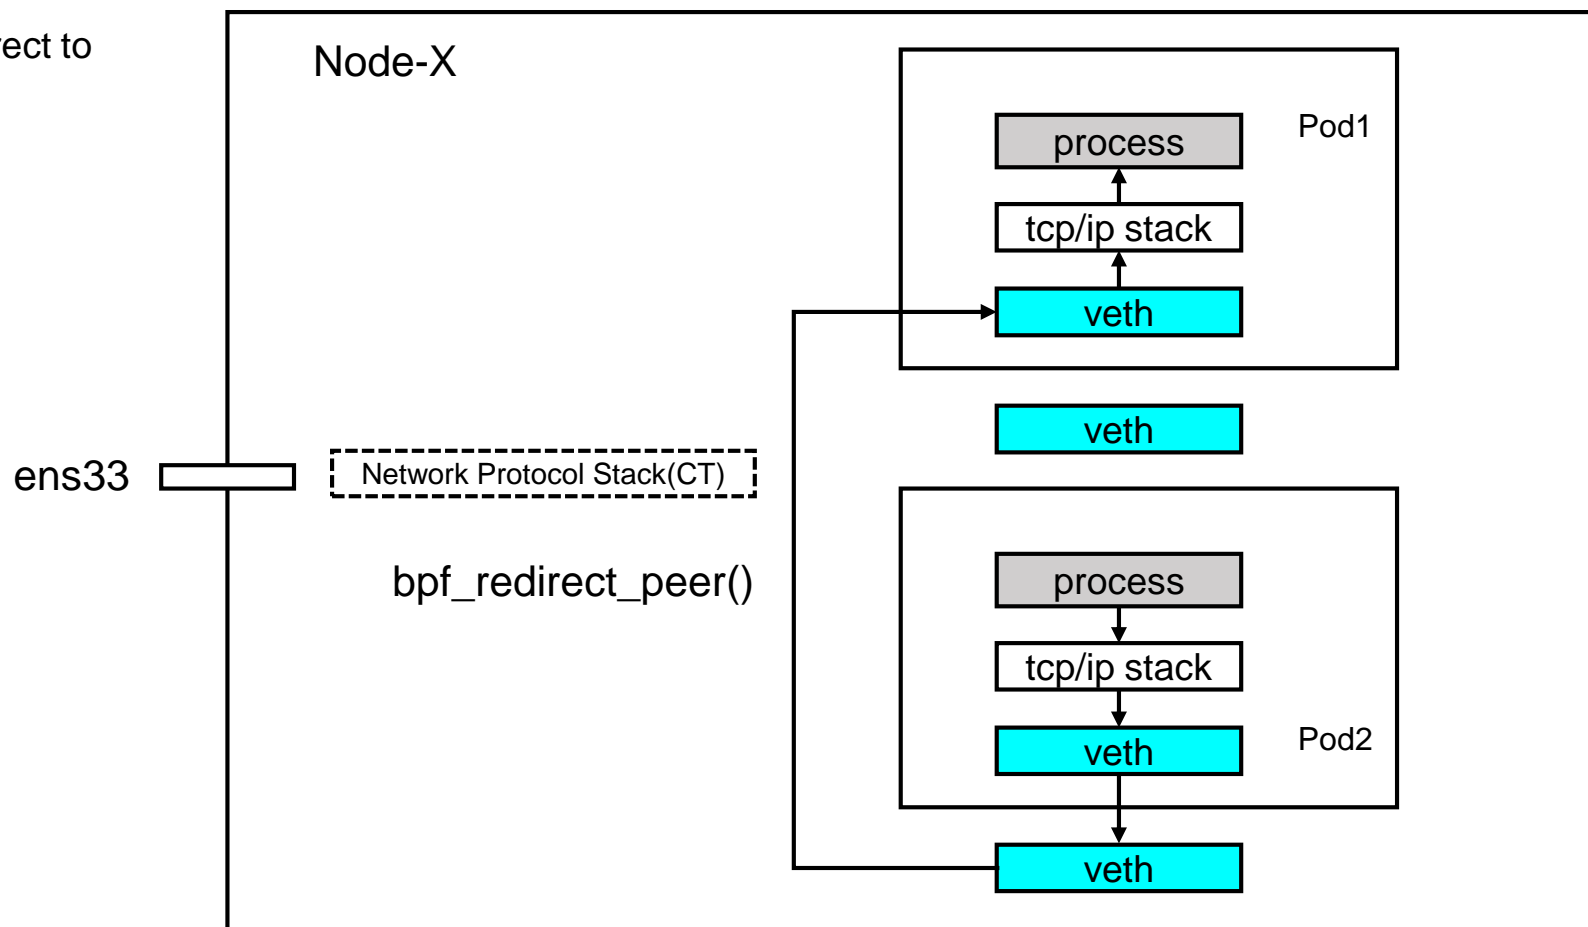


Kubernetes CNI - SRIOV/DPDK Mode



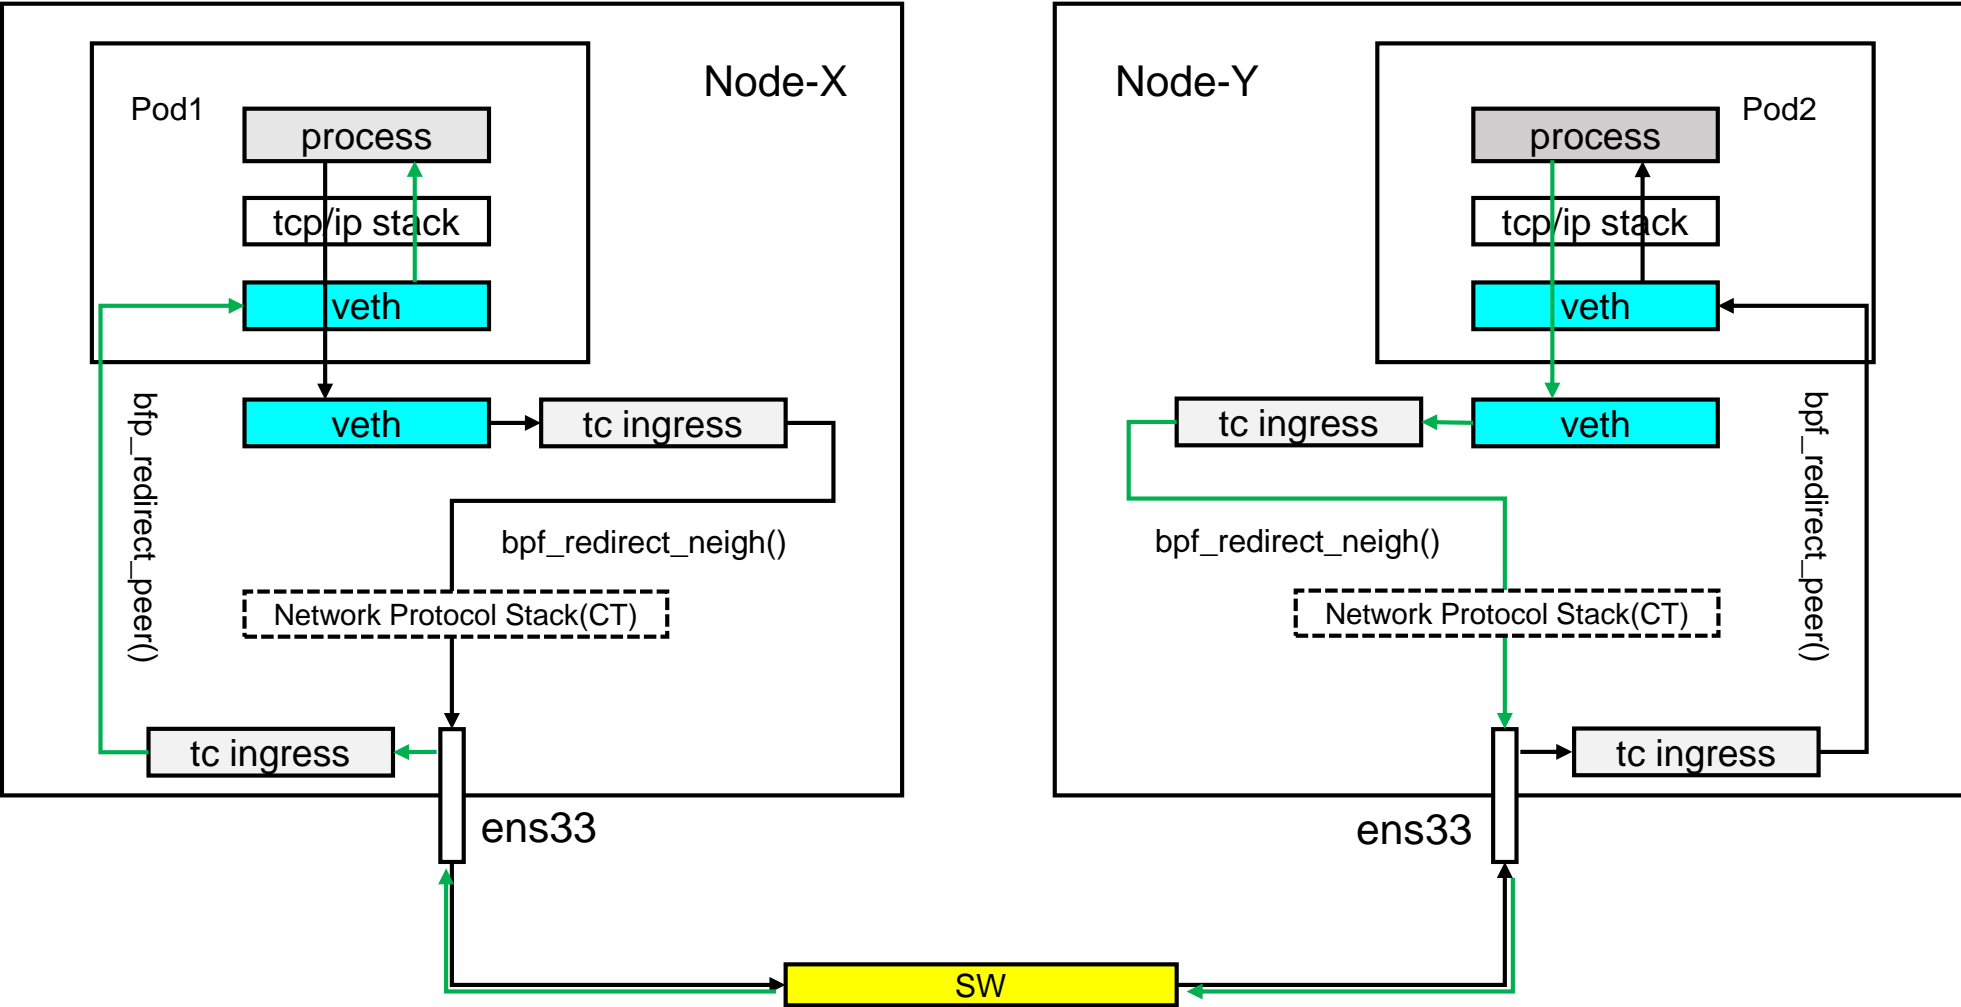
Kubernetes CNI - Cilium eBPF

②: pod namespace 1 redirect to pod namespace 2



The `bpf_redirect_peer()` enables switching network namespaces from the ingress of the NIC to the ingress of the Pod without a software interrupt rescheduling point when traversing the network namespace. The physical NIC can thus push packets up the stack into the application's socket residing in a different Pod namespace in one go. This also leads to quicker application wake-up for picking up the received data. Similarly, rescheduling points are reduced from 2 to 1 for local Pod-to-Pod communication resulting in better latency there as well.

Kubernetes CNI - Cilium eBPF



Kubernetes Network

Copyright And Confidentiality

All content of this material, including text, images, audio, video, software, programs, tools etc., are collected online.

Visitors may use the content or services provided in this material for personal study, research or appreciation, as well as other non-commercial or non-profit uses, but at the same time comply with the provisions of the Copyright Law and other relevant laws, and may not infringe this information and related the legal rights of the right holder. In addition, any use of any content or service of this material for other purposes requires the written permission of the person and payment.



牧云/罗伟

Network | IaaS | PaaS | ServiceMesh

交流 学习 沉淀 成长 分享

olaf.luo@foxmail.com

<https://www.yuque.com/wei.luo>

<https://youdianzhishi.com/web>

All Rights Reserved.

